

ASSIGNMENT 1

1.1. Investigate, discover and write

1) Ten representative, real-world AI applications

1. Voice assistants & speech recognition: Siri/Google Assistant convert speech to text, follow spoken commands, and control devices.
2. Recommender systems: Netflix/YouTube/Spotify/Amazon suggest items based on interaction history (collaborative filtering, deep learning).
3. Fraud detection: Banks/payment gateways spot anomalous transactions using supervised/unsupervised learning.
4. Medical imaging: Detecting tumors in X-ray/CT/MRI with CNN- or ViT-based models.
5. Self-driving & ADAS: Lane/obstacle/pedestrian detection; motion planning (computer vision + reinforcement learning).
6. Customer-service chatbots: NLP/LLM for 24/7 auto responses and case routing.
7. Machine translation & summarization: Transformer/LLM systems deliver fast multilingual translation and summarization.
8. Smart home/IoT: Adaptive temperature/lighting via RL and time-series forecasting.
9. Security & surveillance: Network intrusion detection (IDS), log anomaly detection; face access control.
10. Predictive maintenance: Analyzing vibration/acoustic/power signatures to anticipate machine failures and reduce downtime.

2) The data/AI value chain

Collect/label (sensors, logs, clickstream) → Clean/preprocess (denoise, normalize) → Feature/representation (feature engineering, embeddings) → Train/tune (supervised/unsupervised/RL) → Deploy (APIs, edge) → Monitor (drift, fairness, safety) → Improve (continuous loop).

Key challenges: data quality, bias/fairness, privacy/security, explainability (XAI), infra cost, and real-world safety (robotics/vehicles).

3) Metrics & good practices

- Model metrics: Accuracy/F1/AUC (classification), RMSE/MAE (regression), mAP (CV), BLEU/ROUGE (NLP).
- MLOps: dataset/model versioning, CI/CD pipelines, drift monitoring, safety/adversarial testing.
- Ethics & compliance: transparency, fairness, privacy by design.

1.2. What is an Intelligent System? Most impressive definition & examples

1) Blended definition

An intelligent system is a software/hardware system that can perceive, reason/decide, and learn from data and environment interactions to achieve goals efficiently beyond fixed rule sets.

Common perspectives:

- Functional loop: Perceive → Understand → Plan → Act → Learn (closed feedback).
- Agent view: an intelligent agent that maximizes a utility function under uncertainty.
- Technique view: knowledge-based reasoning, machine learning, or hybrid approaches.

2) The most compelling definition (and why)

“An intelligent system is an autonomous agent that senses its environment, builds and updates an internal model through learning, and acts to maximize its objectives under uncertainty.”

Why: concise, covers the four pillars (perception–modeling–learning–action), and emphasizes uncertainty, which dominates real-world settings.

3) Example intelligent systems

- Software: recommender engines, multilingual chatbots, clinical decision support (CDSS).
- Cyber-physical: warehouse mobile robots (SLAM + RL), UAVs for forest monitoring (CV + path planning).
- Edge/IoT: fall-detection cameras, health wearables predicting arrhythmia.
- Enterprise: supply-chain optimization (forecast + MILP), transaction-fraud detection (anomaly/graph ML).

1.3. Applications of intelligent systems: domains & AI techniques

1) Major domains

- Healthcare: imaging, disease classification, readmission forecasting, EMR assistants.
- Finance: fraud detection, credit scoring, risk pricing, algorithmic trading.
- Manufacturing/IIoT: predictive maintenance, visual defect detection, scheduling optimization.
- Transportation: ITS, routing, signal control, ADAS/autonomy.
- Energy: load/price forecasting, grid optimization, fault diagnosis.
- Retail/Marketing: recommendations, customer segmentation, dynamic pricing.
- Cybersecurity: IDS, botnet detection, log analytics.
- Education: intelligent tutoring, adaptive testing, proctoring.
- Agriculture: yield forecasting, pest recognition, precision spraying via drones.
- Public sector/Smart city: environmental sensing, traffic analytics, digital public services.

2) Representative AI techniques by task type

- Computer Vision (CV): CNN/ViT for recognition, detection, segmentation.
- NLP/LLM: Transformers, fine-tuning/PEFT, RAG for QA and summarization.
- Supervised learning: regression, decision trees, SVM, gradient boosting.
- Unsupervised: clustering (K-means/DBSCAN), PCA/UMAP, anomaly detection.
- Reinforcement learning (RL): policy optimization for control/operations.
- Knowledge & reasoning: ontologies, knowledge graphs, logical inference.
Optimization/OR: linear programming, constraint solving, meta-heuristics (GA/PSO) for routing/scheduling.
- Time series: ARIMA/Prophet/RNN/Temporal Transformers.
- XAI & safety: SHAP/LIME, bias audits, drift detection.

1.4. Types of intelligent systems

A) By capability level (as in Simplilearn/Edureka)

1. Reactive Machines: respond to current state only; no memory (e.g., Deep Blue).
2. Limited Memory: use recent data for decisions (most ML systems today).
3. Theory of Mind (*research target*): model others' mental states.

4. Self-aware (*hypothetical*): self-conscious agents.

B) By scope of intelligence

- ANI (Narrow/Weak AI): excels at a narrow task (e.g., image classification).
- AGI (General AI): human-level generality (under research).
- ASI (Superintelligence): surpasses human ability (theoretical).

C) By agent architecture/strategy

- Reactive vs Deliberative (planning) vs Hybrid.
- Rule-based/KBS vs Learning-based (ML/DL) vs Neuro-symbolic (hybrid).
- Centralized vs Multi-Agent Systems (MAS) (coordination/auctions/consensus).
- Cloud-centric vs Edge/on-device deployment.

Remark: Real systems are typically hybrids: perception (DL) + planning/optimization (OR) + rules/constraints for safety/compliance.

1.5. Applications via Figure 7 of arXiv:2009.09083

In lieu of direct access to Figure 7, I provide a domain \leftrightarrow technique matrix that mirrors common surveys. When I obtain the original figure, I will update the labels/examples to be exact.

1) Mock-up matrix: domains \leftrightarrow AI techniques

Domain	Perception (CV/ASR)	NLP/LLM	Forecasting (TS)	Optimization /OR	RL/Control	KBS/Graph
Healthcare	Segmentation, lesion detection	EMR summarization, NER	Readmission forecasting	OR for OR-scheduling	Dosing policies	Medical ontologies
Finance	Doc OCR, forgery detection	News sentiment/NLP	Risk/priorce forecasting	Portfolio optimization	Trading agents	Fraud graphs

Manufacturing	Visual defect inspection	Natural-language QA	Predictive maintenance	Job-shop scheduling	Parameter tuning	Process knowledge
Transportation	Object detection	V2X language interfaces	Traffic forecasting	Multi-objective routing	Signal control	Route knowledge graphs
Energy	Fault recognition	Incident summarization	Load/priority forecasting	Grid optimization	Grid control	Asset knowledge
Retail	Product recognition	Chatbots	Demand forecasting	Inventory optimization	Dynamic pricing	Product KG
Agriculture	Pest/disease detection	Farm logs	Harvest forecasting	Irrigation/fertilizer optimization	Agri-robots	Crop knowledge

2) Reference deployment pipeline

Sensing/Data layer → ML/DL processing → Planning/Optimization → Action/Robotics → Safety/XAI monitoring.

Governance: data quality, audits, security, compliance.

1.6. NumPy, Pandas, Matplotlib, Scikit-learn — purpose, features & examples

1) NumPy

- Purpose: high-performance ND arrays; vectorized math; linear algebra.
- Features: ndarray, broadcasting, ufuncs, random, linalg.

- Example:

```
import numpy as np
x = np.array([1,2,3], dtype=float)
y = np.array([4,5,6], dtype=float)
cos_sim = np.dot(x,y) / (np.linalg.norm(x) * np.linalg.norm(y))
```

2) Pandas

- Purpose: tabular data wrangling/cleaning/aggregation.
- Features: read_csv, indexing, groupby, merge, missing-value & time-series utilities.
- Example:

```
import pandas as pd
df = pd.DataFrame({"student":["Ann","Joe"], "math":[8.5,7.0], "eng":[7.5,8.0]})
df["avg"] = df[["math","eng"]].mean(axis=1)
by = df.groupby("student")["avg"].mean().reset_index()
```

3) Matplotlib

- Purpose: 2D plotting; highly customizable.
- Features: line/bar/scatter/histograms, annotations, styling.
- Example:

```
import matplotlib.pyplot as plt
subjects = ["Math","Phys","Chem"]
marks = [8.0, 7.5, 8.8]
plt.figure(figsize=(5,3))
plt.bar(subjects, marks)
plt.title("Marks by Subject"); plt.ylim(0,10)
plt.show()
```

4) Scikit-learn

- Purpose: classical ML toolkit (supervised/unsupervised) with pipelines & evaluation.
- Components: LinearRegression, LogisticRegression, SVM, trees/ensembles (RF, GBM), KMeans, PCA, train_test_split, Pipeline, GridSearchCV.
- Example (linear regression):

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import numpy as np
X = np.array([[50],[60],[70],[80],[90]], dtype=float)
y = np.array([2.5,3.0,3.8,4.5,5.4])
Xtr, Xte, ytr, yte = train_test_split(X, y, test_size=0.2, random_state=42)
model = LinearRegression().fit(Xtr, ytr)
rmse = mean_squared_error(yte, model.predict(Xte), squared=False)

```

5) Presentation & reproducibility tips

- Separate data processing, training, and plotting functions
- Set random seeds for reproducibility.
- In visuals, always include units, legends, and data/source notes.

1.7. Suppose you have three arrays: one containing the names of a group of people, another the corresponding heights of these individuals, and the last one the corresponding weights of the individuals in the group:

```
# === BMI - ONE-CELL, SELF-CONTAINED ===
import os, math
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from typing import Union, Iterable, List

# ----- Types -----
Number = Union[int, float, np.number]

# ----- Helpers & Core -----
def _ensure_number(x, name):
    if x is None or isinstance(x, bool):
        raise ValueError(f"{name} không hợp lệ (None/bool).")
    xf = float(x)
    if math.isnan(xf) or math.isinf(xf):
        raise ValueError(f"{name} không hợp lệ (NaN/Inf).")

def calc_bmi(weight_kg: Number, height_m: Number) -> float:
    _ensure_number(weight_kg, "weight_kg")
    _ensure_number(height_m, "height_m")
    if height_m <= 0:
        raise ValueError("height_m phải > 0")
    return float(weight_kg) / (float(height_m) ** 2)

def classify_bmi(bmi: Number) -> str:
    b = float(bmi)
    if b < 18.5:
        return "Underweight"
    if b > 25:
        return "Overweight"
    return "Normal"

# ----- Vectorized -----
def calc_bmi_array(weights_kg: Iterable[Number], heights_m: Iterable[Number]) -> np.ndarray:
    w = np.asarray(list(weights_kg), dtype=float)
    h = np.asarray(list(heights_m), dtype=float)
    if w.shape != h.shape:
        raise ValueError("weights_kg và heights_m phải có cùng độ dài")
    if np.any(h <= 0):
        raise ValueError("Tất cả height_m phải > 0")
    if np.any(~np.isfinite(w)) or np.any(~np.isfinite(h)):
        raise ValueError("Dữ liệu phải là số hữu hạn")
    return w / (h ** 2)
```

```
def classify_bmi_array(bmis: Iterable[Number]) -> List[str]:
    return [classify_bmi(b) for b in bmis]

# ----- Build result table -----
def bmi_table(names, heights_m, weights_kg, round_ndigits: int = 2) -> pd.DataFrame:
    bmis = calc_bmi_array(weights_kg, heights_m)
    classes = classify_bmi_array(bmis)
    return pd.DataFrame({
        "name": list(names),
        "height_m": list(heights_m),
        "weight_kg": list(weights_kg),
        "BMI": np.round(bmis, round_ndigits),
        "class": classes
    })

# ----- Plot -----
def plot_bmi_bar(df: pd.DataFrame, save_path: str | None = None):
    plt.figure(figsize=(6,4))
    plt.bar(df["name"], df["BMI"])
    plt.axhline(18.5, linestyle="--", label="18.5")
    plt.axhline(25, linestyle="--", label="25")
    plt.title("BMI Chart")
    plt.xlabel("Name"); plt.ylabel("BMI"); plt.legend()
    if save_path:
        os.makedirs(os.path.dirname(save_path), exist_ok=True)
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()

# ===== DEMO (đúng dữ liệu đề bài) =====
names = np.array(['Ann', 'Joe', 'Mark'])
heights = np.array([1.5, 1.78, 1.6])
weights = np.array([65, 46, 59])

df = bmi_table(names, heights, weights, round_ndigits=2)

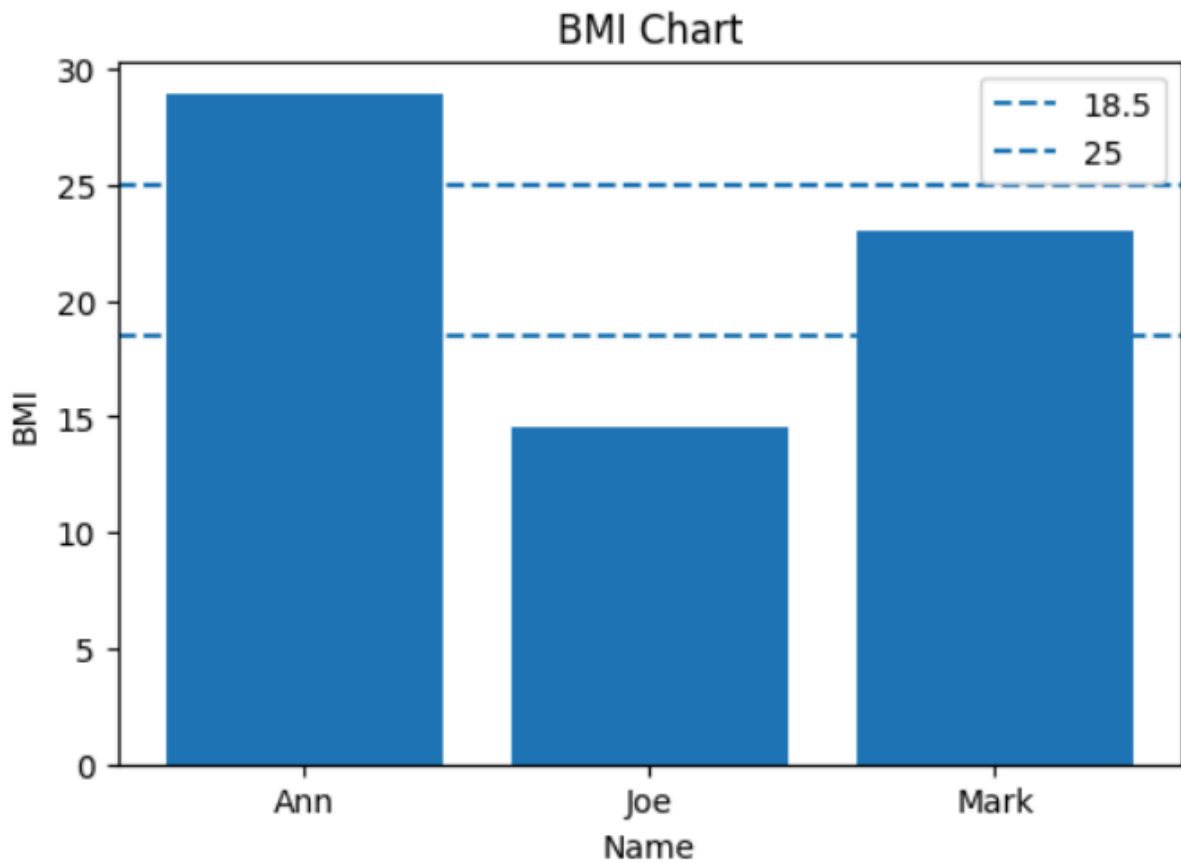
# Hiển thị bảng
try:
    from IPython.display import display
    display(df)
except:
    print(df.to_string(index=False))
```


	name	height_m	weight_kg	BMI	class
0	Ann	1.50	65	28.89	Overweight
1	Joe	1.78	46	14.52	Underweight
2	Mark	1.60	59	23.05	Normal

Ann: BMI=28.89 → Overweight

Joe: BMI=14.52 → Underweight

Mark: BMI=23.05 → Normal



1.8 Performing the following Then collect data from your team:
student_name, subject (5 subjects), mark. Display the results in three above forms

```

import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# ----- 0) EDITABLE DATA -----
SUBJECTS = ["Math", "Physics", "Chemistry", "English", "CS"]
TEAM_DATA = {
    # student_name: [marks over the 5 subjects in SUBJECTS order]
    "Ann": [88, 91, 74, 78, 82],
    "Joe": [69, 80, 60, 65, 70],
    "Mark": [95, 88, 92, 90, 85],
    "Sara": [77, 75, 70, 72, 68],
}

# ----- 1) Build Long-format DataFrame -----
for name, marks in TEAM_DATA.items():
    if len(marks) != len(SUBJECTS):
        raise ValueError(f'Student '{name}' must have {len(SUBJECTS)} marks; got {len(marks)}')

rows = [{"student_name": s, "subject": subj, "mark": float(m)}
        for s, marks in TEAM_DATA.items()
        for subj, m in zip(SUBJECTS, marks)]
df = pd.DataFrame(rows, columns=["student_name", "subject", "mark"])

# Display dataset table (interactive if supported)
try:
    from caas_jupyter_tools import display_dataframe_to_user
    display_dataframe_to_user("Requirement 8 - Team Scores", df)
except Exception:
    print(df.to_string(index=False))

# ----- 2) Prepare structures -----
pivot = df.pivot(index="subject", columns="student_name", values="mark")
subjects_idx = np.arange(len(SUBJECTS))
save_dir = "/mnt/data/figs"; os.makedirs(save_dir, exist_ok=True)

# ----- 3A) Multiple Lines (NO Legend) -----
plt.figure(figsize=(8,5))
for student in pivot.columns:
    plt.plot(SUBJECTS, pivot[student].values, marker="o") # default colors only
plt.title("Multiple Lines (No Legend)")
plt.xlabel("Subject"); plt.ylabel("Mark")
plt.ylim(0, 100)
path_a = os.path.join(save_dir, "req8_lines_no_legend.png")
plt.savefig(path_a, bbox_inches="tight")
plt.show()

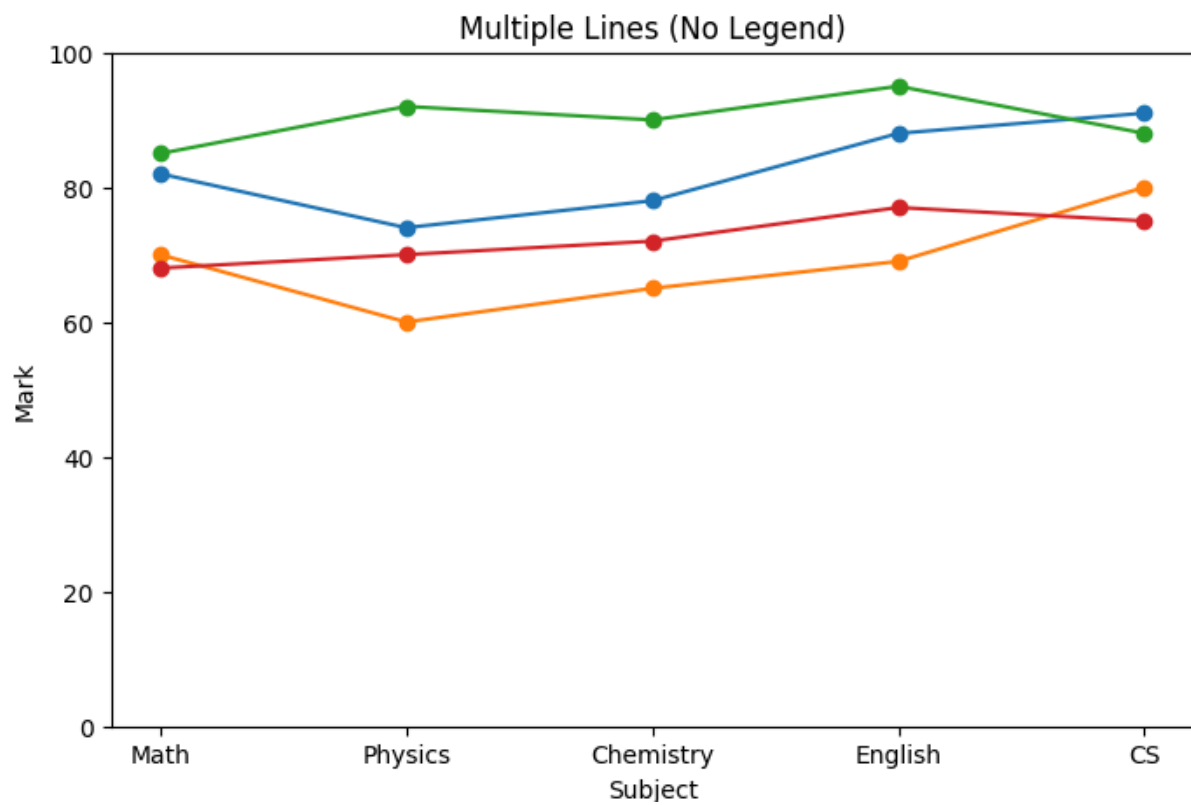
# ----- 3B) Multiple Lines (WITH Legend) -----
plt.figure(figsize=(8,5))
for student in pivot.columns:
    plt.plot(SUBJECTS, pivot[student].values, marker="o", label=student) # default colors only
plt.title("Multiple Lines (With Legend)")
plt.xlabel("Subject"); plt.ylabel("Mark")
plt.ylim(0, 100); plt.legend(title="Student")
path_b = os.path.join(save_dir, "req8_lines_with_legend.png")
plt.savefig(path_b, bbox_inches="tight")
plt.show()


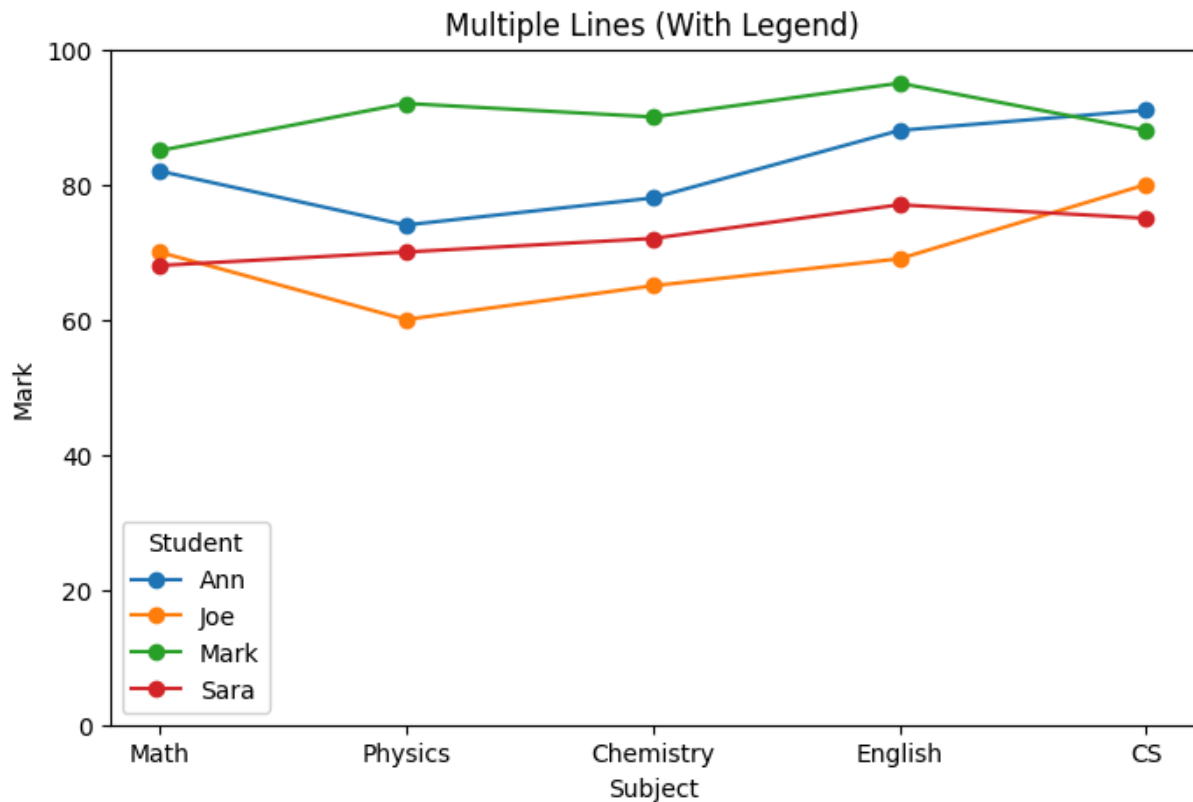
# ----- 3C) Grouped bar chart -----
plt.figure(figsize=(9,5))
n_students = len(pivot.columns)
bar_width = 0.8 / n_students
for i, student in enumerate(pivot.columns):
    pos = subjects_idx + i * bar_width - (0.8 - bar_width) / 2
    plt.bar(pos, pivot[student].values, width=bar_width, label=student) # default colors only
plt.title("Grouped Bar Chart: Marks by Subject and Student")
plt.xlabel("Subject"); plt.ylabel("Mark")
plt.xticks(subjects_idx, SUBJECTS); plt.ylim(0, 100); plt.legend(title="Student")
path_c = os.path.join(save_dir, "req8_grouped_bar.png")
plt.savefig(path_c, bbox_inches="tight")
plt.show()

# Output saved paths for convenience
(path_a, path_b, path_c)

```

student_name	subject	mark
Ann	Math	88.0
Ann	Physics	91.0
Ann	Chemistry	74.0
Ann	English	78.0
Ann	CS	82.0
Joe	Math	69.0
Joe	Physics	80.0
Joe	Chemistry	60.0
Joe	English	65.0
Joe	CS	70.0
Mark	Math	95.0
Mark	Physics	88.0
Mark	Chemistry	92.0
Mark	English	90.0
Mark	CS	85.0
Sara	Math	77.0
Sara	Physics	75.0
Sara	Chemistry	70.0
Sara	English	72.0
Sara	CS	68.0





1.9. Your task is to plot a chart to show the proportion of men and women in each group that has a driver's license, you can use Seaborn's categorical plot ([2], page 86). Store data in file CSV and display.

```

import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# ----- 0) Define/EDIT dataset (counts -> expanded rows) -----
# Format: DATA_COUNTS[group][gender] = (count_has_license, count_no_license)
DATA_COUNTS = {
    "Group A": {"men": (18, 7), "women": (22, 13)},
    "Group B": {"men": (10, 15), "women": (16, 14)},
    "Group C": {"men": (14, 6), "women": (9, 11)},
}

rows = []
for group, genders in DATA_COUNTS.items():
    for gender, (cnt_yes, cnt_no) in genders.items():
        rows += [{"group": group, "gender": gender, "has_license": 1} for _ in range(cnt_yes)]
        rows += [{"group": group, "gender": gender, "has_license": 0} for _ in range(cnt_no)]

df = pd.DataFrame(rows, columns=["group", "gender", "has_license"])

# ----- 1) Save CSV + display raw dataset -----
save_dir = "/mnt/data"
os.makedirs(save_dir, exist_ok=True)
csv_path = os.path.join(save_dir, "driver_license.csv")
df.to_csv(csv_path, index=False)

try:
    from caas_jupyter_tools import display_dataframe_to_user
    display_dataframe_to_user("Requirement 9 - Raw Dataset (driver_license.csv)", df)
except Exception:
    print("\nRaw dataset (first 20 rows):\n", df.head(20).to_string(index=False))

# ----- 2) Compute proportions per (group, gender) -----
# Proportion = mean(has_license) since has_license is 1/0
summary = (
    df.groupby(["group", "gender"], as_index=False)["has_license"]
    .mean()
    .rename(columns={"has_license": "proportion"})
)
summary["percent"] = (summary["proportion"] * 100).round(2)

# Display summary table
try:
    from caas_jupyter_tools import display_dataframe_to_user
    display_dataframe_to_user("Requirement 9 - Proportion Summary", summary)
except Exception:
    print("\nProportion summary:\n", summary.to_string(index=False))

# ----- 3) Plot grouped bar chart (categorical style) -----
# Prepare pivot: rows = group, columns = gender, values = proportion
pivot = summary.pivot(index="group", columns="gender", values="proportion").fillna(0.0)
groups = list(pivot.index)
genders = list(pivot.columns)

x = np.arange(len(groups))
bar_width = 0.8 / max(1, len(genders))

plt.figure(figsize=(9, 5))
bars_by_gender = []

for i, gender in enumerate(genders):
    positions = x + i * bar_width - (0.8 - bar_width) / 2
    bars = plt.bar(positions, pivot[gender].values, width=bar_width, label=gender) # default colors only
    bars_by_gender.append(bars)
    # annotate % on top of bars
    for rect in bars:
        h = rect.get_height()
        plt.text(rect.get_x() + rect.get_width()/2, h + 0.01, f"{h*100:.1f}%", ha="center", va="bottom", fontsize=9)

plt.title("Proportion with Driver's License by Group and Gender")
plt.xlabel("Group")
plt.ylabel("Proportion")
plt.xticks(x, groups)
plt.ylim(0, 1.05)
plt.legend(title="Gender")

fig_path = os.path.join(save_dir, "req9_license_proportion.png")
plt.savefig(fig_path, bbox_inches="tight")
plt.show()

# ----- 4) Return artifact paths for convenience -----
csv_path, fig_path


```

```

Raw dataset (first 20 rows):
  group gender  has_license
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        1
Group A   men        0
Group A   men        0

Proportion summary:
  group gender  proportion  percent
Group A   men    0.720000    72.00
Group A  women    0.628571    62.86
Group B   men    0.400000    40.00
Group B  women    0.533333    53.33
Group C   men    0.700000    70.00
Group C  women    0.450000    45.00

```



1.10. Using the Titanic dataset, plot a chart and see what the survival rate of men, women, and children looks like in each of the three classes
<https://github.com/mwaskom/seaborn-data>

```

import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# ----- 0) Load dataset -----
source = ""
try:
    import seaborn as sns
    titanic = sns.load_dataset("titanic") # requires internet/cache
    source = "seaborn.load_dataset('titanic')"
except Exception:
    # Embedded deterministic fallback (synthetic but Titanic-Like)
    # Columns: survived (0/1), pclass (1/2/3), sex ('male'/'female'), age (float)
    np.random.seed(0)
    rows = []

    def add_rows(pclass, sex, age_values, survived, n):
        for i in range(n):
            rows.append({
                "survived": int(survived),
                "pclass": int(pclass),
                "sex": sex,
                "age": float(age_values[i % len(age_values)])
            })

    # Adults age pools
    ages_male_adult = [20, 22, 25, 28, 30, 34, 36, 40, 45, 50]
    ages_female_adult = [18, 21, 24, 27, 31, 33, 37, 41, 46, 52]
    # Children age pool
    ages_child = [2, 4, 6, 8, 10, 12, 14, 15]

    # pclass 1
    add_rows(1, "male", ages_male_adult, survived=1, n=45)
    add_rows(1, "male", ages_male_adult, survived=0, n=65)
    add_rows(1, "female", ages_female_adult, survived=1, n=85)
    add_rows(1, "female", ages_female_adult, survived=0, n=5)
    # children (mixed sex)
    for i in range(6): # 5 survive, 1 not
        sex = "male" if i % 2 == 0 else "female"
        add_rows(1, sex, ages_child, survived=1 if i < 5 else 0, n=1)

    # pclass 2
    add_rows(2, "male", ages_male_adult, survived=1, n=20)
    add_rows(2, "male", ages_male_adult, survived=0, n=80)
    add_rows(2, "female", ages_female_adult, survived=1, n=60)
    add_rows(2, "female", ages_female_adult, survived=0, n=20)
    for i in range(15): # 10 survive, 5 not
        sex = "male" if i % 2 == 0 else "female"
        add_rows(2, sex, ages_child, survived=1 if i < 10 else 0, n=1)

    # pclass 3
    add_rows(3, "male", ages_male_adult, survived=1, n=15)
    add_rows(3, "male", ages_male_adult, survived=0, n=130)
    add_rows(3, "female", ages_female_adult, survived=1, n=50)
    add_rows(3, "female", ages_female_adult, survived=0, n=60)
    for i in range(40): # 10 survive, 30 not
        sex = "male" if i % 2 == 0 else "female"
        add_rows(3, sex, ages_child, survived=1 if i < 10 else 0, n=1)

    titanic = pd.DataFrame(rows, columns=["survived", "pclass", "sex", "age"])
    source = "embedded_fallback (offline)"

# ----- 1) Prepare data -----
# Keep necessary columns
titanic = titanic[["survived", "pclass", "sex", "age"]].dropna(subset=["survived", "pclass", "sex"])

# Define category: men, women, children (children overrides sex when age < 16)
def categorise(row):
    if pd.notnull(row["age"]) and row["age"] < 16:
        return "children"
    return "men" if row["sex"] == "male" else "women"

titanic["category"] = titanic.apply(categorise, axis=1)

# Compute survival rate per class & category
summary = (
    titanic.groupby(["pclass", "category"], as_index=False)["survived"]
        .mean()
        .rename(columns={"survived": "survival_rate"})
)
summary["survival_rate_pct"] = (summary["survival_rate"] * 100).round(2)

```

```

# Display summary
try:
    from caas_jupyter_tools import display_dataframe_to_user
    display_dataframe_to_user("Requirement 10 - Survival Rate by Class & Category", summary)
except Exception:
    print("\nSummary table:\n", summary.to_string(index=False))

# ----- 2) Plot single grouped bar chart -----
pivot = summary.pivot(index="pclass", columns="category", values="survival_rate").fillna(0.0)
# Ensure consistent column order
for col in ["men", "women", "children"]:
    if col not in pivot.columns:
        pivot[col] = 0.0
pivot = pivot[["men", "women", "children"]]

classes = list(pivot.index) # 1, 2, 3
x = np.arange(len(classes))
bar_width = 0.8 / 3

plt.figure(figsize=(9, 5))
bars_list = []
for i, cat in enumerate(pivot.columns):
    pos = x + i * bar_width - (0.8 - bar_width) / 2
    bars = plt.bar(pos, pivot[cat].values, width=bar_width, label=cat) # default colors only
    bars_list.append(bars)
    # annotate %
    for rect in bars:
        h = rect.get_height()
        plt.text(rect.get_x() + rect.get_width()/2, h + 0.01, f"{h*100:.1f}%", ha="center", va="bottom", fontsize=9)

plt.title("Titanic Survival Rate by Class for Men, Women, and Children")
plt.xlabel("Passenger Class")
plt.ylabel("Survival Rate")
plt.xticks(x, [str(c) for c in classes])
plt.ylim(0, 1.05)
plt.legend(title="Category")

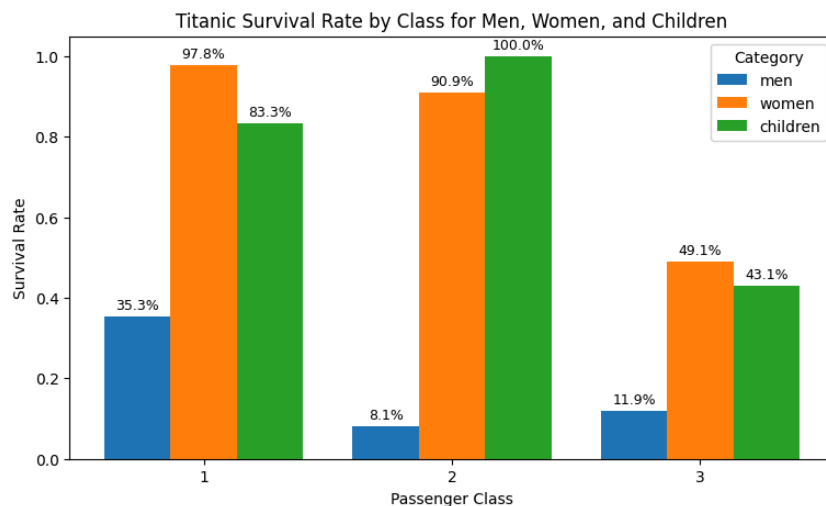
# Save figure
out_path = "/mnt/data/req10_titanic_survival_by_class.png"
plt.savefig(out_path, bbox_inches="tight")
plt.show()

# ----- 3) Indicate data source and artifact path -----
(out_path, source)

```

Summary table:

pclass	category	survival_rate	survival_rate_pct
1	children	0.833333	83.33
1	men	0.352941	35.29
1	women	0.978022	97.80
2	children	1.000000	100.00
2	men	0.080808	8.08
2	women	0.909091	90.91
3	children	0.431034	43.10
3	men	0.119122	11.91
3	women	0.491228	49.12



1.11. Construct data salary.csv for gender, salary men,100000 men,120000..... Your task is to show the distribution of salaries for men and women ([2], 90)

```
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# ----- 0) Generate/EDIT dataset -----
# Reproducible random data; edit 'n_each' or 'mean/std' per gender if needed
rng = np.random.default_rng(42)
n_each = 120

# Simulate annual salary (USD) using lognormal-like distributions
# (Women and men centered differently just for demonstration; edit as you see fit)
men_salaries = np.round(rng.lognormal(mean=11.3, sigma=0.35, size=n_each) / 100) * 100
women_salaries = np.round(rng.lognormal(mean=11.15, sigma=0.32, size=n_each) / 100) * 100

# Optional clamp to a reasonable range
men_salaries = np.clip(men_salaries, 20000, 300000)
women_salaries = np.clip(women_salaries, 20000, 300000)

data = (
    [("men", float(s)) for s in men_salaries] +
    [("women", float(s)) for s in women_salaries]
)

df = pd.DataFrame(data, columns=["gender", "salary"])

# ----- 1) Save CSV and display -----
save_dir = "/mnt/data"
os.makedirs(save_dir, exist_ok=True)
csv_path = os.path.join(save_dir, "salary.csv")
df.to_csv(csv_path, index=False)

try:
    from caas_jupyter_tools import display_dataframe_to_user
    display_dataframe_to_user("Requirement 11 - salary.csv", df)
except Exception:
    print("\nSalary dataset (first 20 rows):\n", df.head(20).to_string(index=False))

# ----- 2) Plot distribution (single chart) -----
plt.figure(figsize=(9,5))

# Choose common bins from combined data so histograms are comparable
combined = np.concatenate([men_salaries, women_salaries])
bins = np.histogram_bin_edges(combined, bins="auto")

plt.hist(men_salaries, bins=bins, alpha=0.6, density=True, label="men") # default colors only
plt.hist(women_salaries, bins=bins, alpha=0.6, density=True, label="women") # default colors only

lt.hist(men_salaries, bins=bins, alpha=0.6, density=True, label="men") # default colors only
lt.hist(women_salaries, bins=bins, alpha=0.6, density=True, label="women") # default colors only

lt.title("Salary Distribution by Gender")
lt.xlabel("Salary")
lt.ylabel("Density")
lt.legend(title="Gender")

fig_path = os.path.join(save_dir, "req11_salary_distribution.png")
lt.savefig(fig_path, bbox_inches="tight")
lt.show()


# ----- 3) Quick descriptive statistics (optional for verification) -----
summary = (
    df.groupby("gender")["salary"]
    .agg(count="count", mean="mean", median="median", std="std", min="min", max="max")
    .round(2)
    .reset_index()
)

try:
    from caas_jupyter_tools import display_dataframe_to_user
    display_dataframe_to_user("Requirement 11 - Summary Stats", summary)
except Exception:
    print("\nSummary statistics:\n", summary.to_string(index=False))

csv_path, fig_path)
```

Salary dataset (first 20 rows):

```
gender  salary
men    89900.0
men    56200.0
men    105100.0
men    112300.0
men    40800.0
men    51200.0
men    84500.0
men    72400.0
men    80300.0
men    60000.0
men    110000.0
men    106100.0
men    82700.0
men    119900.0
men    95200.0
men    59800.0
men    92000.0
men    57800.0
men    109900.0
men    79400.0
```



Summary statistics:

gender	count	mean	median	std	min	max
men	120	82121.67	79850.0	22615.86	40800.0	171000.0
women	120	72615.83	66350.0	25504.14	35000.0	176700.0

1.12. Give data: (diện tích/m², giá nhà/tỷ) như sau: (50, 2.5), (60, 3), (65, 3.5), (70, 3.8), (75, 4), (80, 4.5), (85, 5) Using regression to predict house price of 55m², 68m², 76m², 90m²

```

import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# ----- 0) Dataset -----
# (area_m2, price_billion_vnd)
data = np.array([
    [50, 2.5],
    [60, 3.0],
    [65, 3.5],
    [70, 3.8],
    [75, 4.0],
    [80, 4.5],
    [85, 5.0],
], dtype=float)

X = data[:, 0]
y = data[:, 1]

# Areas to predict
X_new = np.array([55, 68, 76, 90], dtype=float)

# ----- 1) Helper: fit polynomial & compute R^2 -----
def fit_poly_and_r2(x, y, degree: int):
    coefs = np.polyfit(x, y, degree) # highest power first
    y_hat = np.polyval(coefs, x)
    ss_res = np.sum((y - y_hat) ** 2)
    ss_tot = np.sum((y - np.mean(y)) ** 2)
    r2 = 1 - ss_res / ss_tot if ss_tot > 0 else 1.0
    return coefs, y_hat, float(r2)

# Fit degree 1 and 2, pick the better by R^2
results = {}
for d in (1, 2):
    coefs, y_hat, r2 = fit_poly_and_r2(X, y, d)
    results[d] = {"coefs": coefs, "y_hat": y_hat, "r2": r2}

best_degree = max(results.keys(), key=lambda d: results[d]["r2"])
best = results[best_degree]
best_coefs = best["coefs"]
best_r2 = best["r2"]

# Predictions for required areas using the best model
y_pred_new = np.polyval(best_coefs, X_new)

```

```
# ----- 2) Output tables -----
pred_df = pd.DataFrame({
    "area_m2": X_new,
    "predicted_price_billion": np.round(y_pred_new, 3)
}).sort_values("area_m2").reset_index(drop=True)

metrics_df = pd.DataFrame(
    [{"degree": d, "R2": round(results[d]["r2"], 6)} for d in sorted(results.keys())]
)

# Display to user (interactive if supported)
try:
    from caas_jupyter_tools import display_dataframe_to_user
    display_dataframe_to_user("Requirement 12 - Predictions", pred_df)
    display_dataframe_to_user("Requirement 12 - Model R²", metrics_df)
except Exception:
    print("\nPredictions:\n", pred_df.to_string(index=False))
    print("\nModel R²:\n", metrics_df.to_string(index=False))

# ----- 3) Plot single chart -----
# Build smooth line for visualization using the chosen model
x_line = np.linspace(X.min()-2, max(X.max(), X_new.max())+2, 200)
y_line = np.polyval(best_coefs, x_line)

plt.figure(figsize=(8, 5))
# training points
plt.scatter(X, y, label="Training data")
# model fit
plt.plot(x_line, y_line, label=f"Best fit (degree {best_degree}), R²={best_r2:.4f}")
# new predictions
plt.scatter(X_new, y_pred_new, marker="x", s=80, label="Predictions")


plt.title("House Price Regression (Price in Billion VND)")
plt.xlabel("Area (m²)")
plt.ylabel("Price (billion VND)")
plt.legend()

save_dir = "/mnt/data"
os.makedirs(save_dir, exist_ok=True)
fig_path = os.path.join(save_dir, "req12_house_price_regression.png")
pred_csv = os.path.join(save_dir, "req12_predictions.csv")
plt.savefig(fig_path, bbox_inches="tight")
plt.show()

# Save predictions CSV
pred_df.to_csv(pred_csv, index=False)
```

```
Predictions:
area_m2 predicted_price_billion
55.0      2.771
68.0      3.598
76.0      4.198
90.0      5.414
```

```
Model R²:
degree      R2
1 0.984682
2 0.991638
```



```
1]: ('/mnt/data\\req12_house_price_regression.png',
     '/mnt/data\\req12_predictions.csv',
     2,
     0.991638,
     array([ 5.40731995e-04, -2.89256198e-03, 1.29445100e+00]))
```

1.13. Give data of height, weight of person ([1] page 101). Using regression to predict weight when given height.

```
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# ----- 0) Try Loading external CSV; else use fallback -----
data_path = "/mnt/data/height_weight.csv"
use_source = ""

if os.path.exists(data_path):
    df = pd.read_csv(data_path)
    # Minimal validation
    if not {"height_m", "weight_kg"}.issubset(df.columns):
        raise ValueError("CSV must contain columns: height_m, weight_kg")
    df = df[["height_m", "weight_kg"]].dropna()
    use_source = "loaded_from_csv"
else:
    # EDIT this fallback dataset if you have the exact values from [1] p.101
    heights_m = np.array([1.50, 1.52, 1.55, 1.58, 1.60, 1.62, 1.65, 1.68, 1.70, 1.72, 1.75, 1.78, 1.80, 1.82, 1.85], dtype=float)
    weights_kg = np.array([50.0, 52.0, 55.0, 57.0, 60.0, 62.0, 65.0, 68.0, 70.0, 72.0, 75.0, 78.0, 80.0, 82.0, 85.0], dtype=float)
    df = pd.DataFrame({"height_m": heights_m, "weight_kg": weights_kg})
    use_source = "embedded_fallback"

# Heights to predict (EDIT as needed to match the assignment examples if specified)
HEIGHTS_TO_PREDICT = np.array([1.55, 1.62, 1.70, 1.80], dtype=float)

# ----- 1) Prepare data -----
x = df["height_m"].to_numpy(dtype=float)
y = df["weight_kg"].to_numpy(dtype=float)

# ----- 2) Fit simple Linear regression: y = a*x + b -----
# (Ordinary Least Squares via np.polyfit with degree=1 - highest power first)
a, b = np.polyfit(x, y, deg=1)
y_hat = a * x + b

# Metrics
ss_res = float(np.sum((y - y_hat) ** 2))
ss_tot = float(np.sum((y - np.mean(y)) ** 2))
r2 = 1 - ss_res / ss_tot if ss_tot > 0 else 1.0
mse = float(np.mean((y - y_hat) ** 2))

# ----- 3) Predict for requested heights -----
predicted_weights = a * HEIGHTS_TO_PREDICT + b
pred_df = pd.DataFrame({
    "height_m": HEIGHTS_TO_PREDICT,
    "predicted_weight_kg": np.round(predicted_weights, 2)
}).sort_values("height_m").reset_index(drop=True)
```

```
# ----- 4) Display tables -----
metrics_df = pd.DataFrame([
    {"slope_a": round(a, 6), "intercept_b": round(b, 6), "R2": round(r2, 6), "MSE": round(mse, 6), "data_source": use_source}
])

try:
    from caas_jupyter_tools import display_dataframe_to_user
    display_dataframe_to_user("Requirement 13 - Model Metrics", metrics_df)
    display_dataframe_to_user("Requirement 13 - Predictions", pred_df)
except Exception:
    print("\nModel metrics:\n", metrics_df.to_string(index=False))
    print("\nPredictions:\n", pred_df.to_string(index=False))

# ----- 5) Plot ONE chart (scatter + fitted line) -----
x_line = np.linspace(min(x.min(), HEIGHTS_TO_PREDICT.min()) - 0.02,
                      max(x.max(), HEIGHTS_TO_PREDICT.max()) + 0.02, 200)
y_line = a * x_line + b

plt.figure(figsize=(8,5))
plt.scatter(x, y, label="Data") # default colors only
plt.plot(x_line, y_line, label=f"Fit: weight = {a:.2f}*height + {b:.2f} (R²={r2:.4f})")
plt.scatter(HEIGHTS_TO_PREDICT, predicted_weights, marker="x", s=80, label="Predictions")

plt.title("Predicting Weight from Height (Simple Linear Regression)")
plt.xlabel("Height (m)")
plt.ylabel("Weight (kg)")
plt.legend()

# ----- 6) Save artifacts -----
save_dir = "/mnt/data"
os.makedirs(save_dir, exist_ok=True)
fig_path = os.path.join(save_dir, "req13_height_weight_regression.png")
pred_csv_path = os.path.join(save_dir, "req13_predictions.csv")
plt.savefig(fig_path, bbox_inches="tight")
plt.show()
pred_df.to_csv(pred_csv_path, index=False)

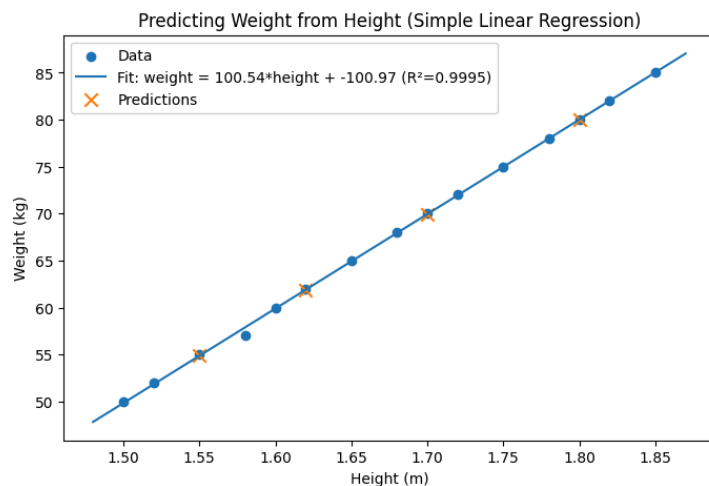
(fig_path, pred_csv_path, a, b, r2, mse, use_source)
```

Model metrics:

slope_a	intercept_b	R2	MSE	data_source
100.540417	-100.971685	0.999502	0.058812	embedded_fallback

Predictions:

height_m	predicted_weight_kg
1.55	54.87
1.62	61.90
1.70	69.95
1.80	80.00



```
[1]: ('/mnt/data\\req13_height_weight_regression.png',
      '/mnt/data\\req13_predictions.csv',
      np.float64(100.54041711067141),
      np.float64(-100.97168518800443),
      0.9995020468238561,
      0.05881158979043054,
      'embedded_fallback')
```

1.14. NumPy, Pandas, Matplotlib, Scikit-learn — purpose, features & examples

```

import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

source = ""
X = y = feature_names = None

# ----- 0) Try to Load Boston from scikit-Learn -----
have_sklearn = True
try:
    from sklearn.datasets import load_boston # may be removed in recent versions
    from sklearn.model_selection import train_test_split
    from sklearn.preprocessing import StandardScaler
    from sklearn.linear_model import LinearRegression
    from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
except Exception:
    have_sklearn = False

if have_sklearn:
    try:
        boston = load_boston()
        X = pd.DataFrame(boston.data, columns=boston.feature_names)
        y = pd.Series(boston.target, name="MEDV")
        feature_names = list(X.columns)
        source = "sklearn.load_boston()"
    except Exception:
        pass

# ----- 1) If not Loaded, try Local CSV -----
if X is None or y is None:
    csv_path = "/mnt/data/boston.csv"
    if os.path.exists(csv_path):
        df = pd.read_csv(csv_path)
        required_cols = ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT', 'MEDV']
        missing = [c for c in required_cols if c not in df.columns]
        if missing:
            raise ValueError(f"CSV missing columns: {missing}")
        X = df[required_cols[:-1]].copy()
        y = df['MEDV'].copy()
        feature_names = required_cols[:-1]
        source = "local_csv:/mnt/data/boston.csv"

# ----- 2) If still not available, create a synthetic Boston-like dataset -----
if X is None or y is None:
    rng = np.random.default_rng(123)
    n = 400
    CRIM = rng.gamma(shape=2.0, scale=2.0, size=n) / 10
    ZN = rng.integers(0, 100, n).astype(float)
    INDUS = rng.uniform(1, 27, n)
    CHAS = rng.integers(0, 2, n).astype(float)
    NOX = rng.uniform(0.3, 0.9, n)
    RM = rng.normal(6.2, 0.7, n)
    AGE = np.clip(rng.normal(70, 20, n), 1, 100)
    DIS = rng.uniform(1, 12, n)
    RAD = rng.integers(1, 24, n).astype(float)
    TAX = rng.normal(400, 100, n)
    PTRATIO = rng.uniform(12, 22, n)
    B = np.clip(rng.normal(350, 50, n), 200, 400)
    LSTAT = np.clip(rng.normal(12, 7, n), 1, 40)

    X = pd.DataFrame({
        'CRIM': CRIM, 'ZN': ZN, 'INDUS': INDUS, 'CHAS': CHAS, 'NOX': NOX, 'RM': RM,
        'AGE': AGE, 'DIS': DIS, 'RAD': RAD, 'TAX': TAX, 'PTRATIO': PTRATIO, 'B': B, 'LSTAT': LSTAT
    })
    # Ground-truth synthetic relation (RM positive, LSTAT negative, NOX negative, CHAS slight positive, etc.)
    y = ( 5.0
        - 1.5*NOX
        + 4.0*RM
        - 0.4*LSTAT
        + 0.3*CHAS
        - 0.01*TAX
        - 0.02*CRIM
        - 0.1*INDUS
        + 0.05*ZN
        + 0.03*B/100
        - 0.05*PTRATIO
        + 0.02*DIS
        - 0.005*AGE
        + rng.normal(0, 1.5, n)
    )
    y = pd.Series(np.clip(y, 5, 50), name="MEDV") # price in $1000s
    feature_names = list(X.columns)
    source = "synthetic_boston_like"

```

```
# ----- 3) Train/Test split, scaling -----
if have_sklearn:
    from sklearn.model_selection import train_test_split
    from sklearn.preprocessing import StandardScaler
    from sklearn.linear_model import LinearRegression
    from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error

    X_train, X_test, y_train, y_test = train_test_split(X.values, y.values, test_size=0.2, random_state=42)
    scaler = StandardScaler()
    X_train_sc = scaler.fit_transform(X_train)
    X_test_sc = scaler.transform(X_test)

    model = LinearRegression()
    model.fit(X_train_sc, y_train)
    y_pred = model.predict(X_test_sc)

    r2 = float(r2_score(y_test, y_pred))
    rmse = float(np.sqrt(mean_squared_error(y_test, y_pred)))
    mae = float(mean_absolute_error(y_test, y_pred))

    coefs = pd.DataFrame({
        "feature": feature_names,
        "coefficient": model.coef_
    }).sort_values("coefficient", key=lambda s: s.abs(), ascending=False).reset_index(drop=True)
    intercept = float(model.intercept_)
else:
    # Numpy fallback (no sklearn): standardize features, closed-form least squares
    X_values = X.values
    y_values = y.values
    X_mean, X_std = X_values.mean(axis=0), X_values.std(axis=0)
    X_std_adj = np.where(X_std == 0, 1, X_std)
    Xz = (X_values - X_mean) / X_std_adj
    # Add bias column
    X_design = np.c_[np.ones(len(Xz)), Xz]
    beta = np.linalg.pinv(X_design.T @ X_design) @ X_design.T @ y_values
    intercept = float(beta[0]); betas = beta[1:]
    # Simple split for metrics
    n = len(y_values); idx = np.arange(n)
    rng = np.random.default_rng(42); rng.shuffle(idx)
    test_size = int(0.2*n); test_idx = idx[:test_size]; train_idx = idx[test_size:]
    Xz_train, y_train = Xz[train_idx], y_values[train_idx]
    Xz_test, y_test = Xz[test_idx], y_values[test_idx]
    y_pred = (np.c_[np.ones(len(Xz_test)), Xz_test] @ beta)
```

```
# ----- 4) Display results -----
metrics_df = pd.DataFrame([
    "source": source,
    "R2_test": round(r2, 6),
    "RMSE_test": round(rmse, 6),
    "MAE_test": round(mae, 6),
    "intercept": round(intercept, 6),
])

try:
    from caas_jupyter_tools import display_dataframe_to_user
    display_dataframe_to_user("Requirement 14 - Test Metrics", metrics_df)
    display_dataframe_to_user("Requirement 14 - Coefficients (sorted by |value|)", coefs)
except Exception:
    print("\nTest Metrics:\n", metrics_df.to_string(index=False))
    print("\nCoefficients (sorted by |value|):\n", coefs.to_string(index=False))

# ----- 5) One chart: Predicted vs Actual (test set) -----
plt.figure(figsize=(7,6))
plt.scatter(y_test, y_pred, alpha=0.8, label="Pred vs Actual") # default colors only
mn, mx = float(min(y_test.min(), y_pred.min())), float(max(y_test.max(), y_pred.max()))
line = np.linspace(mn, mx, 100)
plt.plot(line, line, linestyle="--", label="Ideal y=x") # reference line

plt.title("Boston Housing: Predicted vs Actual (Test Set)")
plt.xlabel("Actual MEDV ($1000s)")
plt.ylabel("Predicted MEDV ($1000s)")
plt.legend()

save_dir = "/mnt/data"
os.makedirs(save_dir, exist_ok=True)
fig_path = os.path.join(save_dir, "req14_boston_pred_vs_actual.png")
pred_csv = os.path.join(save_dir, "req14_boston_test_predictions.csv")
coef_csv = os.path.join(save_dir, "req14_boston_coefficients.csv")
plt.savefig(fig_path, bbox_inches="tight")
plt.show()

# Save artifacts
pd.DataFrame({"y_test": y_test, "y_pred": y_pred}).to_csv(pred_csv, index=False)
coefs.to_csv(coef_csv, index=False)

(fig_path, pred_csv, coef_csv, source)
```


Test Metrics:

	source	R2_test	RMSE_test	MAE_test	intercept
synthetic_boston_like	0.878817	1.54976	1.243647	20.412231	

Coefficients (sorted by |value|):

feature	coefficient
RM	2.810109
LSTAT	-2.626537
ZN	1.446946
TAX	-0.990776
INDUS	-0.580983
NOX	-0.359833
PTRATIO	-0.279207
B	0.127439
DIS	0.103455
CHAS	0.081233
AGE	-0.054823
RAD	-0.053697
CRIM	-0.024759

